

Problem 1.

For this question, you will compute word counts in a large corpus of text. Word frequency lists (a frequency of the word is the ratio of the number of times the word appears to the total number of words in a document) are commonly used in computational linguistics. You will write a function that finds the 10 most frequently occurring words in a text file, and a program that uses this function to find the 10 most frequently-occurring words in *Pride and Prejudice*. For the purposes of this problem, you may assume, if you like, that all words are separated by spaces and that there is no punctuation in the file, and all the words are in lowercase, so that the list of words in the file `text.txt` is given by

```
open("text.txt", encoding="latin-1").read().split()
```

Part (a)

First, store the number of times that the word `w` appears in the text in `word_counts[w]`. For example, if the word “the” appears in the file 5 times, `word_counts["the"]` should be 5.

For example, if the contents of your file are the opening of *Notes from the Underground* by Fyodor Dostoyevsky, translated by Constance Garnett:

I am a sick man. I am a spiteful man. I am an unattractive man. I believe my liver is diseased.
However, I know nothing at all about my disease, and do not know for certain what ails me.

, and you do not process the text in any way other than using `read().split()`, the list of words will be:

```
["I", "am", "a", "sick", "man.", "I", "am", "a", "spiteful", "man.", "I", "am",  
"an", "unattractive", "man.", "I", "believe", "my", "liver", "is", "diseased.",  
"However,", "I", "know", "nothing", "at", "all", "about", "my", "disease,", "and",  
"do", "not", "know", "for", "certain", "what", "ails", "me."]
```

`word_counts` should then be

```
{"sick": 1, "man.": 3, "at": 1, "what": 1, "nothing": 1, "do": 1, "is": 1, "me.": 1,  
"I": 5, "ails": 1, "an": 1, "am": 3, "know": 2, "disease,": 1, "not": 1, "liver": 1,  
"believe": 1, "all": 1, "my": 2, "certain": 1, "However,": 1, "and": 1, "for": 1,  
"unattractive": 1, "spiteful": 1, "about": 1, "a": 2, "diseased.": 1}
```

For example, the word `"man."` (with the period at the end) appears 3 times in the text, so its entry in the dictionary `word_counts` is 3.

Part (b)

Write a function with the signature `top10(L)` that takes in a list `L` of 100 different integers, and returns a list of the 10 largest integers in `L`.

Part (c)

Now, obtain the top 10 most-frequent words from the dictionary `freq`. To do that, you need to sort the data by the word counts. You cannot sort dictionaries directly, but you can use the following trick:

```
inv_freq = {6: "the", 12: "a", 1:"hi"}  
print(sorted(inv_freq.items()))
```

First, experiment with this code and understand what it is doing, and then apply the technique to finding the top 10 most frequent words. Test your function by creating a small text file where you can find the top n most-frequent words manually, running your code on this file, and comparing the results. Then, download the text of *Pride and Prejudice* from <http://www.gutenberg.org/files/1342/1342-0.txt> and obtain the top 10 most frequent words in *Pride and Prejudice*.

Problem 2.

Most web pages are written in HTML (HyperText Markup Language). HTML is a fairly complicated language, but here are some basics:

- Paragraphs are enclosed between `<p>` and `</p>`
- Bolded text is enclosed between `` and ``
- A link to <http://www.google.com>, which appears as the text “Google” would look like this:

```
<a href = "http://www.google.com/">Google</a>
```

Download `hello.html` from <http://www.cs.toronto.edu/~guerzhoy/180/labs/hello.html>. To do that, open the link in a web browser, and save the HTML file by pressing [Cntr-s] and specifying the location to which to save the file in the dialogue box that pops up. Now open your local copy of `hello.html` that you just saved with the text editor. On Windows, use the Notepad text editor. On Mac, use TextEdit. Use the Plain Text Mode in TextEdit <https://apple.stackexchange.com/questions/17433/can-textedit-save-as-plain-text>.

Also open the local copy of `hello.html` in a web browser. To view the source of the HTML file, click inside the page in the browser with the right mouse button/cntrl-click on a Mac, and select “View source”.

Modify `hello.html` as follows.

- Make the word “ever” appear in bold
- Add a link to Yahoo.ca’s search results for “engineering science.”

Show the results to a TA, in the browser and the text editor. By looking at your answer to the previous question, figure out a way to search Yahoo by modifying the text in the address text box instead of typing search terms in the search text box in the browser. Note that when searching Yahoo, the address line in the browser changes depending on the search terms.

Problem 3.

Yahoo.com displays the number of results for the search terms entered on the bottom left of the results page. Use [Right Click]→[View Page Source] to see the HTML source of the results page, and find

in the HTML file the number of results is (to do that, search for the number of the search results that you see in Firefox in the HTML file). Now, write a function in Python that takes in a search term and returns the number of results for that search term on Yahoo.com. To do that, you need to automatically download the Yahoo! search results page. Files on the web can be read similarly to how you would read local files. For example, you can read and then display the source code of a file on my webpage using:

```
import urllib.request
f = urllib.request.urlopen("http://www.cs.toronto.edu/~guerzhoy/180/draft/draft.html")
page = f.read().decode("utf-8")
f.close()
print(page)
```

Here's an excerpt from Dave Barry's Ask Mister Language Person column:

Q. Do you take questions from attorneys? A. Yes. That will be \$ 475. Q. No, seriously, Im an attorney, and I want to know which is correct: "With regards to the aforementioned blah blah blah." Or: "With regards to the aforementioned yak yak yak." A. That will be \$ 850.

http://articles.baltimoresun.com/1991-10-13/features/1991286206_1_blah-yak-yak-transpire

And here's an excerpt from the *New York Times* column *On Language*:

Neale Gifford writes: **"One practice that annoys me is the use of 'one-year anniversary' or 'five year anniversary' instead of 'first' or 'fifth.' Reason? Anniversary is derived from the Latin annus meaning 'year.'" Ed Morman writes of "n-year anniversary," "It's not very mellifluous and it is, of course redundant. Help me bring back 'nth anniversary.'"**

<http://www.nytimes.com/2010/07/18/magazine/18onlanguage-anniversary.html>

One way to settle those usage disputes is to check which variant appears on the web more often¹

Write a function with the signature `choose_variant(variants)` which takes in a list of usage variants `variants` and returns the variant which returns the most search results on Yahoo.com. For example,

`choose_variant(["five-year anniversary", "fifth anniversary"])` should return

"fifth anniversary", since that phrase appears more often. (When searching for a phrase, we enclose it in quotes to indicate that we're searching for a phrase and not for a set of keywords.) Now call `choose_variant()` with `["top ranked school uoft", "top ranked school waterloo"]`.

¹While this method is used by writers as well as linguists to determine which linguistic constructs are more popular, the counts are actually not quite accurate. See, e.g., here: <http://itre.cis.upenn.edu/~myl/languageblog/archives/001834.html>